

Database Programming with SQL

9-1: Using GROUP BY and HAVING Clauses

Practice Activities

Objectives

- Construct and execute a SQL query using GROUP BY
- Construct and execute a SQL query using GROUP BY ... HAVING
- Construct and execute a GROUP BY on more than one column
- Nest group functions

Vocabulary

Identify the vocabulary word for each definition below.

- Used to specify which groups are to be displayed; restricts groups that do not meet group criteria **HAVING**
- Divides the rows in a table into groups **GROUP BY**

Try It / Solve It

1. In the SQL query shown below, which of the following is true about this query?

___ **T** ___ a. Kimberly Grant would not appear in the results set.

___ **F** ___ b. The GROUP BY clause has an error because the manager_id is not listed in the SELECT clause.

___ **F** ___ c. Only salaries greater than 16001 will be in the result set.

___ **F** ___ d. Names beginning with Ki will appear after names beginning with Ko.

___ **F** ___ e. Last names such as King and Kochhar will be returned even if they don't have salaries > 16000.

```
SELECT last_name, MAX(salary)
FROM employees
WHERE last_name LIKE 'K%'
GROUP BY manager_id, last_name
HAVING MAX(salary) > 16000
```

```
ORDER BY last_name DESC ;
```

2. Each of the following SQL queries has an error. Find the error and correct it. Use Oracle Application Express to verify that your corrections produce the desired results.

a. SELECT manager_id
FROM employees
WHERE AVG(salary) < 16000

```
GROUP BY manager_id;  
SELECT manager_id  
FROM employees  
GROUP BY manager_id  
HAVING AVG(salary) < 16000;
```

b.

```
SELECT cd_number, COUNT(title)  
FROM d_cds  
WHERE cd_number < 93;  
SELECT cd_number, COUNT(title)  
FROM d_cds  
WHERE cd_number < 93  
GROUP BY cd_number;
```

c.

```
SELECT ID, MAX(ID), artist AS Artist  
FROM d_songs  
WHERE duration IN('3 min', '6 min', '10 min')  
HAVING ID < 50  
GROUP by ID;  
SELECT ID, artist AS Artist  
FROM d_songs  
WHERE duration IN ('3 min', '6 min', '10 min')  
AND ID < 50;
```

d.

```
SELECT loc_type, rental_fee AS Fee  
FROM d_venues  
WHERE id < 100  
GROUP BY "Fee"  
ORDER BY 2;  
SELECT loc_type, rental_fee AS Fee  
FROM d_venues  
WHERE id < 100  
ORDER BY 2;
```

3. Rewrite the following query to accomplish the same result:

```
SELECT DISTINCT MAX(song_id)  
FROM d_track_listings  
WHERE track IN ( 1, 2, 3);  
SELECT MAX(song_id)  
FROM d_track_listings  
WHERE track IN (1, 2, 3)  
GROUP BY track;
```

4. Indicate True or False

___T___ a. If you include a group function and any other individual columns in a SELECT clause, then each individual column must also appear in the GROUP BY clause.

___F___ b. You can use a column alias in the GROUP BY clause.

___F___ c. The GROUP BY clause always includes a group function.

5. Write a query that will return both the maximum and minimum average salary grouped by department from the employees table.

```
SELECT MAX(avg_salary) AS max_avg_salary, MIN(avg_salary) AS min_avg_salary
FROM (
    SELECT department_id, AVG(salary) AS avg_salary
    FROM employees
    GROUP BY department_id
) avg_salaries;
```

6. Write a query that will return the average of the maximum salaries in each department for the employees table.

```
SELECT AVG(max_salary) AS avg_max_salary
FROM (
    SELECT department_id, MAX(salary) AS max_salary
    FROM employees
    GROUP BY department_id
) dept_max_salaries;
```

Database Programming with SQL

9-2: Using ROLLUP and CUBE Operations and GROUPING SETS

Practice Activities

Objectives

- Use ROLLUP to produce subtotal values
- Use CUBE to produce cross-tabulation values
- Use GROUPING SETS to produce a single result set
- Use the GROUPING function to identify the extra row values created by either a ROLLUP or CUBE operation

Vocabulary

Identify the vocabulary word for each definition below.

- Used to create subtotals that roll up from the most detailed level to a grand total, following a grouping list specified in the clause **ROLLUP**
- An extension to the GROUP BY clause like ROLLUP that produces cross-tabulation reports **CUBE**
- Used to specify multiple groupings of data **GROUP BY**

Try It / Solve It

1. Within the Employees table, each manager_id is the manager of one or more employees who each have a job_id and earn a salary. For each manager, what is the total salary earned by all of the employees within each job_id? Write a query to display the

Manager_id, job_id, and total salary. Include in the result the subtotal salary for each manager and a grand total of all salaries.

```
SELECT manager_id, job_id,  
       SUM(salary) AS total_salary  
FROM employees  
GROUP BY ROLLUP(manager_id, job_id)  
ORDER BY manager_id, job_id;
```

2. Amend the previous query to also include a subtotal salary for each job_id regardless of the manager_id.

```
SELECT manager_id, job_id,  
       SUM(salary) AS total_salary  
FROM employees  
GROUP BY ROLLUP(job_id, manager_id)  
ORDER BY job_id, manager_id;
```

3. Using GROUPING SETS, write a query to show the following groupings:

- department_id, manager_id, job_id
- manager_id, job_id
- department_id, manager_id

```
SELECT department_id, manager_id, job_id,  
       SUM(salary) AS total_salary  
FROM employees  
GROUP BY  
       GROUPING SETS (  
           (department_id, manager_id, job_id), (manager_id, job_id), manager_id, job_id  
           (department_id, manager_id) )  
ORDER BY department_id, manager_id, job_id;
```

Database Programming with SQL

9-3: Set Operators

Practice Activities

Objectives

- Define and explain the purpose of SET operators
- Use a set operator to combine multiple queries into a single query
- Control the order of rows returned using set operators

Vocabulary

Identify the vocabulary word for each definition below.

- operator that returns all rows from both tables and eliminates duplicates **UNION**
- columns that were made up to match queries in another table that are not in both tables
Derived Columns
- operator that returns all rows from both tables, including duplicates **UNION ALL**
- used to combine results into one single result from multiple SELECT statements **Set operators**

- operator that returns rows that are unique to each table **MINUS**
- operator that returns rows common to both tables **INTERSECT**

Try It / Solve It

1. Name the different Set operators?

UNION:

```
SELECT column_name(s) FROM table1
UNION
SELECT column_name(s) FROM table2;
```

UNION ALL:

```
SELECT column_name(s) FROM table1
UNION ALL
SELECT column_name(s) FROM table2;
```

INTERSECT:

```
SELECT column_name(s) FROM table1
INTERSECT
SELECT column_name(s) FROM table2;
```

EXCEPT:

```
SELECT column_name(s) FROM table1
EXCEPT (or MINUS)
SELECT column_name(s) FROM table2;
```

2. Write one query to return the employee_id, job_id, hire_date, and department_id of all employees and a second query listing employee_id, job_id, start_date, and department_id from the job_history table and combine the results as one single output. Make sure you suppress duplicates in the output.

```
SELECT employee_id, job_id, hire_date AS start_date, department_id
FROM employees
UNION
SELECT employee_id, job_id, start_date, department_id
FROM job_history;
```

3. Amend the previous statement to not suppress duplicates and examine the output. How many extra rows did you get returned and which were they? Sort the output by employee_id to make it easier to spot.

```
SELECT employee_id, job_id, hire_date AS start_date, department_id
FROM employees
UNION ALL
SELECT employee_id, job_id, start_date, department_id
FROM job_history
ORDER BY employee_id;
```

4. List all employees who have not changed jobs even once. (Such employees are not found in the job_history table)

```
SELECT employee_id, job_id, hire_date, department_id
FROM employees
```

WHERE employee_id NOT IN (SELECT employee_id FROM job_history);

5. List the employees that HAVE changed their jobs at least once.

SELECT employee_id, job_id, hire_date, department_id

FROM employees

INNER JOIN job_history

ON employee_id = employee_id;

6. Using the UNION operator, write a query that displays the employee_id, job_id, and salary of ALL present and past employees. If a salary is not found, then just display a 0 (zero) in its place.

SELECT employee_id, job_id, salary

FROM employees

UNION ALL

SELECT employee_id, job_id, COALESCE(salary, 0) AS salary

FROM job_history;